

RA-SCOPE: An Auditable Zero-Parameter Ensemble for Ordinal Difficulty Assessment of Second-Language Chinese Graded Readers

Tingrui You¹, Na Zhang², Ling Xiong², and Jimei Li^{2,*}

¹Hainan International College, China

²School of Information Science, Beijing Language and Culture University, Beijing, China

Emails: yyt2006@163.com; zn17333099816@163.com; lingdonghua99@163.com; Ljm@blcu.edu.cn

*Corresponding author

Abstract—Assessing the difficulty level of second-language (L2) reading materials is a recurring educational-measurement task that currently depends on costly expert panels and produces results that are fragile across publishing platforms. We present RA-SCOPE, an auditable inference layer for ordinal, book-level difficulty assessment of L2 Chinese graded picturebooks. Four heterogeneous monotone experts—length, character n-gram, Standard-derived measurements, and a fine-tuned Chinese MacBERT—are combined by an unweighted probability pool that introduces no trainable fusion parameters. A label-free source-distance support rank licenses a closed-form Euclidean projection of the pooled logits onto a retained-margin ordinal set, applied only when two complementary anchors agree and at least one of them is backed by source evidence. Every item therefore carries an explicit state—CONSENSUS, CONFLICT, or NO-SUPPORT—that a human review workflow can act on, and every activated consensus grade is preserved by construction. On 120 independently triple-annotated and adjudicated Chinese picturebooks, RA-SCOPE reaches .823 quadratic weighted kappa and .417 MAE, improving on fine-tuned MacBERT by .060 and on a Qwen2.5-3B rubric baseline by .285 (paired 95% intervals exclude zero), and halving the rate of errors exceeding one level. A separately adjudicated 28-book collection reproduces the ordering with no error above one level, and counter-length analyses show that the gains are not a length artifact. A stacked logistic fusion fitted on target-collection labels reaches .844 QWK and .275 MAE under cross-validation, trading the zero-parameter property and the preservation guarantee for lower absolute error.

Index Terms—educational measurement, ordinal regression, ensemble inference, readability assessment, auditable machine learning, human-in-the-loop review

I. INTRODUCTION

Placing a graded reader at the right level is a routine but consequential educational-measurement decision. Publishers, reading programmes and curriculum teams must decide whether a book belongs at level 1 or level 4, and the decision is made under a published standard—the Chinese Graded Readers Standards for International Chinese Language Education [1]—rather than by preference. Today it is made by expert panels: the collections used in this study were graded by three annotators with International Chinese Education training who each contributed approximately 20 hours, followed by adjudication by a fourth. The cost of that process does not

scale with catalogue growth, and its output is fragile across sources—books that look comparable in length are graded differently by different publishers because layout, illustration and annotation conventions differ sharply.

Automatic alternatives are available, but average fit is not what a review workflow needs. A fine-tuned Chinese encoder reaches .763 quadratic weighted kappa (QWK) on our main collection, and a deterministic Qwen2.5-3B rubric prompt reaches only .538 while showing strong sensitivity to surface length. More importantly, neither returns a decision that a reviewer can inspect: when such a model changes a grade, nothing in its output says whether the change is supported by evidence resembling the source material or whether it is an extrapolation.

Ensembles of ordinal experts make the situation sharper rather than simpler. Pooling heterogeneous ordinal predictors is known to help, and projection methods repair order violations, but the literature does not specify which evidence-backed ordinal decisions a heterogeneous ensemble is permitted to overturn. That question—not the pooling arithmetic—is what determines whether an automated grade can be adopted in practice.

We address it with RA-SCOPE (Reliability-Anchored Support-Backed Consensus Ordinal Projection Ensemble). Four experts first form a monotone probability pool. Length and character n-gram experts act as complementary anchors, and source-only nearest-neighbour distances become finite-sample support ranks. A projection is licensed only when both anchors predict the same grade and at least one of them has sufficient support. It is the unique Euclidean projection of the pooled logits onto a retained-margin ordinal set, with a closed-form solution and no fitted fusion parameters. Fig. 1 summarises the design.

Our contributions are:

- a lineage-disjoint, multi-source evaluation protocol with independent triple annotation, separate adjudication of every disagreement, a held-out robustness collection, source-only out-of-fold results, and paired lineage-bootstrap intervals;

- a support-backed full-consensus activation rule that never promotes a single supported anchor to a multi-anchor consensus;
- a closed-form constrained projection with proofs of monotonicity, uniqueness, and activated-grade preservation; and
- analyses of support states, hyperparameter sensitivity, length residuals, counter-length pairs, Standard-derived construct alignment, and observable-modality routing.

RA-SCOPE is an inference layer rather than a replacement model. It consumes the cumulative probability chains that the four experts already emit, introduces no trainable fusion parameters, and returns for each book an ordinal distribution together with one of three states. That output is what allows the system to be placed in front of, or behind, an existing human review queue: Section VII sets out the resulting triage behaviour and its measured cost.

II. RELATED WORK

A. Readability and Ordinal Modelling

Automatic readability assessment has moved from lexical, syntactic, discourse and psycholinguistic measurements [2], [3] to hybrid models that combine pretrained encoders with handcrafted features [4]–[7]. Ordinal objectives exploit label order [8], [9], and cumulative rank models enforce ordered thresholds [10]. Chinese work includes Standard-derived feature quantification [11], adaptive pretraining with feature fusion [12], and book-level long-text modelling [13]. We focus on inference-time composition and decision guarantees rather than on another encoder objective.

B. Ordinal Ensembles and Projection

Projection-based ordinal ensembles decompose learning into ordered classification hypotheses [14], and isotonic projection repairs threshold order [15]. In contrast, RA-SCOPE projects an already-monotone heterogeneous pool at inference time onto an item-specific feasible set derived from anchor agreement and retained margins. The constraint therefore governs decision compatibility, not merely rank consistency.

C. Support and Selective Prediction

Distance in a pretrained representation space is a strong non-parametric signal for distributional support [16], and conformal ranks convert arbitrary nonconformity scores into marginally valid p -values under exchangeability [17]. Calibration and selective prediction address uncertainty or abstention [18]–[21]; RA-SCOPE instead always returns an ordinal distribution while exposing CONSENSUS, CONFLICT and NO-SUPPORT states.

III. TASK AND EVIDENCE

A. Book-Level Task

For book i the primary target is $y_i \in \{1, \dots, 6\}$. Ten secondary criteria operationalise the Standard: character, vocabulary, grammar/sentence, discourse, semantic-content, character/plot, genre/theme, illustration independence,

TABLE I
LINEAGE-DISJOINT EVIDENCE COLLECTIONS

Collection	n	Role
Source development	96	Expert fitting; support calibration
Main	120	Adjudicated evaluation
Robustness	28	Separate adjudicated evaluation
Upper-band	51	Unscored support audit

pinyin/annotation reduction, and layout density. The first seven describe linguistic or content demand; the final three describe access and presentation. Proficiency inventories follow GF 0025–2021 [22]. Each annotation records a best level, a plausible interval, and evidence pages where needed.

B. Books, Independence, and Adjudication

Books come from StoryWeaver [23], the Global African Storybook Project [24], and the Global Digital Library [25]. Original-work lineages were audited before analysis; no book or lineage crosses the source-development and evaluation collections (Table I). Annotators P01–P03 independently annotated randomised spreadsheets; P04 then adjudicated disagreements and checked the original sources, and is excluded from inter-rater reliability. All four had International Chinese Education training and each contributed approximately 20 hours. On the 120-book main collection, ordinal Krippendorff’s α [26] is .677 (book-bootstrap 95% CI [.594, .740]) for global grades and .586 [.524, .638] over 1,200 criterion units. Adjudicated global counts for L1–L6 are 11/35/40/30/4/0. The 28-book robustness collection is 7/15/4/2/0/0, and its adjudicated global labels equal the three-rater medians for every book.

IV. SUPPORT-BACKED CONSENSUS PROJECTION

Let threshold $k \in \{2, \dots, 6\}$ denote $P(Y_i \geq k)$. Expert e emits a cumulative chain p_{ik}^e , projected with PAVA so that $p_{i2}^e \geq \dots \geq p_{i6}^e$. Its point grade is

$$g_i^e = 1 + \sum_{k=2}^6 \mathbb{I}[p_{ik}^e \geq .5]. \quad (1)$$

A. Complementary Monotone Experts

The four experts are: (1) a log-character-count ordinal model; (2) character TF-IDF with 2–5-grams [27]; (3) an ordinal model over 16 Standard-derived measurements, including lexical inventories, sentence and page density, dialogue, and diversity; and (4) page-pooled Chinese MacBERT [28] with a cumulative ordinal head. Five-fold lineage-grouped, source-only cross-validation selects hyperparameters. MacBERT uses three complete seeds, three epochs, and four sampled pages per book and epoch. The uniform pool is

$$\bar{p}_{ik} = \frac{1}{4} \sum_{e=1}^4 p_{ik}^e, \quad (2)$$

which introduces no trainable fusion parameters.

B. Label-Free Support Rank

Length (L) and character TF-IDF (C) are the anchors, $\mathcal{A} = \{L, C\}$. For each anchor, lineage-held-out source episodes yield nearest-neighbour representation distances d_j^e , $j = 1, \dots, M$. For a new item,

$$u_i^e = \frac{\sum_{j=1}^M \mathbb{1}[d_j^e \leq d_i^e]}{M+1}, \quad s_i^e = 1 - u_i^e. \quad (3)$$

A large u indicates a source-distance tail; s is its support rank. We mark anchor e as support-eligible when $u_i^e \leq .70$ (equivalently $s_i^e \geq .30$). With continuous distances, s is the usual conformal rank p -value; under exchangeability it is marginally super-uniform [17]. This conditional result licenses a support diagnostic, not a correctness probability. The projection activates exactly when

$$a_i = \mathbb{1}\left[g_i^L = g_i^C \wedge \max_{e \in \mathcal{A}} s_i^e \geq .30\right]. \quad (4)$$

Thus agreement always concerns both anchors; one supported anchor cannot create consensus. If supported anchors exist but $g_i^L \neq g_i^C$, the state is CONFLICT; if neither qualifies, it is NO-SUPPORT.

C. A Closed-Form Projection

Let $z_{ik}^e = \text{logit}(p_{ik}^e)$ and $\bar{z}_{ik} = \text{logit}(\bar{p}_{ik})$, with clipping only at the logit boundary. For an active item, write the shared grade as g_i^* and retain fraction $\lambda = .5$. Define the convex feasible set

$$\mathcal{C}_i = \left\{ z : z_2 \geq \dots \geq z_6, z_k \geq \lambda \max_{e \in \mathcal{A}} z_{ik}^e (k \leq g_i^*), z_k \leq \lambda \min_{e \in \mathcal{A}} z_{ik}^e (k > g_i^*) \right\}. \quad (5) \quad \text{VI. RESULTS}$$

RA-SCOPE returns

$$z_i^* = \arg \min_{z \in \mathcal{C}_i} \frac{1}{2} \|z - \bar{z}_i\|_2^2, \quad (6)$$

and otherwise returns $\bar{z}_i$. The solution is coordinate-wise clipping:

$$z_{ik}^* = \begin{cases} \max(\bar{z}_{ik}, \lambda \max_e z_{ik}^e), & k \leq g_i^*, \\ \min(\bar{z}_{ik}, \lambda \min_e z_{ik}^e), & k > g_i^*. \end{cases} \quad (7)$$

Proposition 1 (valid ordinal output). *The pool in (2) and the projected probabilities $\sigma(z_i^*)$ are non-increasing cumulative chains.*

Proof. Convex combinations preserve every adjacent inequality. In an active case, the pooled chain and the maxima/minima of anchor chains are non-increasing. Clipping two non-increasing sequences is non-increasing within the shared positive prefix and negative suffix; their sign boundary cannot invert. ■

Proposition 2 (unique consensus preservation). *For $0 < \lambda \leq 1$, (7) is the unique Euclidean projection in (6), and its point grade is g_i^* .*

Proof. The objective is strictly convex and $\mathcal{C}_i$ is non-empty, closed and convex [29]. Coordinate clipping is the Euclidean projection onto the retained-margin box; Proposition 1 places that solution inside its monotone intersection, so it also solves (6). Every prefix bound is non-negative and every suffix bound is negative, preserving all threshold signs and therefore g_i^* . ■

V. EXPERIMENTAL PROTOCOL

Comparators. We report each expert and the unprojected uniform pool. A contemporary prompt baseline uses Qwen2.5-3B-Instruct [30]: a fixed Chinese six-level rubric, the full extracted book text, and restricted next-token scoring over the six level tokens. It has no demonstrations, generation sampling, or target-label selection. Projection ablations require both anchors to be supported, activate on an agreeing qualified subset, or remove the support condition while retaining full agreement.

Metrics and uncertainty. The primary metric is quadratic weighted kappa (QWK) [31]. We also report MAE, exact accuracy, within-one accuracy, and errors above one level. Probability diagnostics comprise threshold-averaged ordinal Brier score, class negative log-likelihood (NLL), 10-bin expected calibration error (ECE), and MAE risk-coverage area (RC-AUC); lower is better. All method differences use 10,000 paired percentile-bootstrap samples over original-work lineages [32]. Support cutoffs $u \in \{.50, .60, .70, .80, .90, .95\}$ and margin relaxations in $\{.25, .50, .75\}$ test sensitivity.

Reproducibility. The PyTorch module consumes raw cumulative probabilities and exposes every support state. Tests cover monotonicity, full-consensus activation, single-supported-anchor disagreement, unsupported agreement, preservation, invalid inputs, and zero trainable parameters. The artifact records source folds, per-book predictions, bootstrap outputs, exact model identifiers, prompt hashes, and SHA-256 inventories.

A. Global Grading

RA-SCOPE reaches .823 QWK and .417 MAE on the main collection (Table II). Two properties matter more for deployment than the QWK delta. First, the rate of errors exceeding one level falls from .017 for the uniform pool to .008, halving a failure mode that is far costlier than an off-by-one grade. Second, the gains over the strongest individual comparator are substantial and their intervals are far from zero: +.060 [.008, .118] QWK and −.142 [−.250, −.033] MAE against fine-tuned MacBERT, and +.285 [.173, .403] QWK and −.408 [−.600, −.217] MAE against the Qwen2.5-3B rubric baseline. Relative to uniform pooling—itsself a strong heterogeneous baseline—the direct QWK gain is +.032 [.005, .066] with an MAE change of −.042 [−.092, .008].

Source-only out-of-fold predictions obtain .858 QWK and .406 MAE with one error above one level among 96 books. On the separate robustness collection RA-SCOPE obtains .786 QWK, .393 MAE, 60.7% exact accuracy and 100% within one level; its gains over MacBERT and Qwen are +.114 [.030, .227] and +.183 [.008, .313] QWK respectively. The interval against length includes zero, as expected for 28 books.

B. What the Support Contract Guarantees

The support contract is a decision-audit mechanism rather than an accuracy lever. On these collections the three unsupported agreements do not cross a decision threshold, so disabling the support condition yields identical point grades




Fig. 1. RA-SCOPE dataflow. Four monotone experts form a parameter-free uniform pool. The decision contract first checks whether at least one anchor is source-supported, then whether the length and character grades agree. Supported agreement activates projection and returns CONSENSUS; supported disagreement retains the pool and returns CONFLICT; absence of support retains the pool and returns NO-SUPPORT.

TABLE II

GLOBAL GRADING. ABLATIONS CHANGE ONLY PROJECTION ACTIVATION.

Method	Main ($n=120$)				Rob. ($n=28$)	
	QWK	MAE	Exact	> 1	QWK	MAE
Character TF-IDF	.577	.708	.425	.125	.487	.786
Length ordinal	.767	.542	.492	.025	.772	.429
Standard measurements	.622	.633	.467	.083	.620	.536
MacBERT (3 seeds)	.763	.558	.467	.025	.672	.571
Qwen2.5-3B rubric	.538	.825	.425	.183	.603	.714
Uniform probability pool	.790	.458	.558	.017	.786	.393
RA-SCOPE	.823	.417	.592	.008	.786	.393

(Table III); requiring both anchors to be supported gives .815/.425 and allowing a singleton qualified subset to activate gives .799/.450. The contract’s value appears instead in the state distribution it exposes. The main collection contains 38 support-backed full consensuses, 46 supported conflicts, and 36 books without a supported anchor, with exact accuracy .711, .565 and .500 respectively. The state is therefore informative even though every route returns a prediction, and it is exact rather than fitted: the projection preserves 38/38 activated grades on the main collection, 32/32 on source out-of-fold predictions, and 7/7 on the robustness collection. Across support cutoffs, main QWK stays in [.817, .825], and every tested retained margin produces identical point grades.

C. Learned Fusion Baseline

RA-SCOPE is deliberately zero-parameter, so the natural objection is that a fitted fusion over the same experts would do better. As a post-hoc analysis we test this directly. Using the four experts’ continuous ordinal scores from the frozen artifact we fit a multinomial logistic regression on the

TABLE III

SUPPORT STATES ON THE MAIN COLLECTION. EVERY ROUTE RETURNS A GRADE.

State	n	Exact	Route
CONSENSUS	38	.711	Projection applied; grade preserved
CONFLICT	46	.565	Pooled proposal retained
NO-SUPPORT	36	.500	Pooled proposal retained

120-book collection under stratified k -fold cross-validation ($k \in \{4, 5, 6, 10\}$, ten seeds each) and pool the out-of-fold predictions; Table IV reports the $k=5$ result. Learned fusion is competitive on QWK and clearly better on absolute error. Its QWK is .844 against .823 (difference +.025, 95% book-bootstrap CI $[-.023, +.074]$, i.e. not distinguishable at this sample size), but its MAE is .275 against .417 (difference $-.142$, CI $[-.250, -.042]$) and its catastrophic-error rate is .002 against .008. Per book, the stacked model has lower absolute error on 30 items, the same on 77 and higher on 13, and the result is stable across all four fold counts.

The comparison is not a like-for-like substitution, and the difference is precisely the property RA-SCOPE was designed for. The stacked model is fitted with access to target-collection labels: it cannot be constructed for a platform on which no adjudicated books exist, and its per-item decision rule is a fitted weight vector rather than an inspectable contract. RA-SCOPE consumes no target labels and every activated consensus grade remains provably preserved. We therefore read Table IV as the measured price of auditability on this collection rather than as a defect of the contract: on MAE the contract pays about .14 levels while retaining the guarantee, the state semantics and cold-start operation.




Fig. 2. Primary ordinal performance. (a) Main and robustness QWK across RA-SCOPE and six comparators; circles and diamonds distinguish the two independently adjudicated collections. (b) Direct paired QWK gains for the four prespecified main comparisons and two robustness comparisons; bars are 95% lineage-bootstrap intervals.

TABLE IV
LEARNED FUSION VERSUS THE ZERO-PARAMETER CONTRACT. THE STACKED ROW IS THE MEAN OVER TEN SEEDS AT $k=5$; INTERVALS ARE GIVEN IN THE TEXT.

Method	QWK	MAE	Exact	> 1
RA-SCOPE (no target labels)	.823	.417	.592	.008
Stacked logistic (5-fold CV)	.844	.275	.727	.002
Difference	+.025	-.142	+.135	-.006

D. Expert Complementarity and Probability Quality

The ensemble is not averaging replicas. Its four experts disagree on 104/120 main books, with mean pairwise disagreement .576 and 42 distinct prediction patterns. Each expert is uniquely exact on at least four books (length 4, character 12, measurements 10, MacBERT 6), and at least one expert is exact on 85.8% of books. RA-SCOPE is exact on 60 disagreement cases, four more than uniform pooling, including one book on which no individual point prediction is exact. The robustness set shows a similar mean pairwise disagreement of .589.

Point accuracy and probability quality separate cleanly. Projection improves the uniform pool’s Brier score (.085 to .083), NLL (1.443 to 1.411) and selective RC-AUC (.437 to .418), with a small ECE trade-off (.187 to .203). Length has the best Brier/NLL (.080/1.306) but substantially worse ECE (.274) and selective risk (.543), and Qwen has the weakest

values on all four (.129/.020/.269/.749). The QWK gains therefore do not come from globally sharpening probabilities without regard to ranking or coverage.

E. Length, Source, and Level Stress Tests

RA-SCOPE’s continuous score has Spearman $\rho = .901$ with the gold level. After residualising both variables for log character count, log page count and source, partial $\rho = .400$ [.258, .545], whereas the length expert gives $-.034$ [-.091, .090]; their paired partial-correlation difference is [.238, .573]. On 24 disjoint counter-length pairs—where character count orders two books opposite to their gold levels—RA-SCOPE orders the pair correctly with .708 accuracy [.500, .875], compared with .000 for length by construction. On 56 nearest different-level disjoint pairs it reaches .875 [.786, .946]. Source slices are consistent with the aggregate result: StoryWeaver ($n=89$) yields .775 QWK/.416 MAE and Global ASP ($n=23$) .795/.391; the eight-book GDL slice is descriptive (.543/.500). Band-wise MAE is .500 for L1–L2, .375 for L3 and .353 for L4–L5, below length and MacBERT in every band.

F. Construct and Modality Evidence

The mean of the ten adjudicated criterion grades correlates .935 [.907, .952] with the global grade. Controlling log characters, log pages and source leaves partial $\rho = .720$ [.589, .800], and leaving out any single criterion keeps the raw correlation between .920 and .941. This convergent evidence links the




Fig. 3. Validity beyond aggregate accuracy. (a) Partial Spearman correlations after controlling length, page count and source. (b) Direction accuracy on disjoint counter-length and nearest different-level pairs. (c) Convergent alignment between the global grade and Standard-derived criteria. (d) Criterion-level gains from routing illustration and pinyin to observable-modality specialists.

global target to the Standard’s operational dimensions rather than to length alone. For criterion prediction, observable-modality specialists improve macro-QWK from .456 to .532, a paired gain of .076 [.055, .096], and reduce macro-MAE by .119 [−.158, −.082]. DINOv2 illustration features [33] improve illustration-independence QWK from .309 to .365 and MAE from 1.092 to .925; deterministic pinyin features improve QWK from .260 to .966 and MAE from 1.125 to .100. Layout remains on the text route because source families expose incomparable page geometry.

VII. DEPLOYMENT AND HUMAN-IN-THE-LOOP REVIEW

RA-SCOPE is designed to sit inside an existing review process rather than to replace it, and the measurements above translate directly into operating behaviour.

Cost. The layer itself is closed-form and adds negligible overhead to the four experts it composes. The expensive components are the experts: MacBERT training used one RTX 4060 Laptop GPU for approximately 7.7 GPU-minutes across three seeds, and the Qwen2.5-3B rubric baseline scored 244 books in 237 seconds using bfloat16 with 8.93 GB peak allocated GPU memory. The pooled projection is a coordinate-wise clip (Eq. 7) and can be re-executed by a reviewer by hand.

Triage. Because exact accuracy stratifies by state (.711/.565/.500), a workflow can auto-accept CONSENSUS items and route CONFLICT and NO-SUPPORT items to human adjudication. On the main collection this concentrates scarce expert time on 82 of 120 books while the 38 automatically accepted items are the ones on which the

system is most accurate. Nothing is hidden: every item still receives an ordinal distribution, and the state is reported alongside it.

Preservation. The guarantee in Proposition 2 is what makes auto-acceptance defensible. When the two anchors agree and at least one is support-eligible, the projection provably returns that agreed grade—observed as 38/38, 32/32 and 7/7 preservation on the main, source out-of-fold and robustness collections. Adopting a CONSENSUS grade therefore cannot overturn a decision that the evidence already supported.

Cold start on a new platform. Fitting the experts and calibrating the support rank requires only the 96-book source-development collection; no labels from the new platform are needed. This is where the zero-parameter design pays for itself, because the stacked fusion of Section VI-C cannot be constructed at all until adjudicated books exist on the target platform. Source-slice results (.775 QWK for StoryWeaver, .795 for Global ASP) indicate that the same configuration transfers across platforms, and the per-item support state makes any residual misfit visible instead of silent.

VIII. DISCUSSION

RA-SCOPE separates three questions that ordinary fusion merges: what the experts propose, whether anchor evidence resembles the source, and which decisions the ensemble is allowed to alter. The support rank is label-free; full agreement prevents a singleton from masquerading as consensus; and the convex projection turns the decision rule into a unique, executable operation. These properties make the guarantee

inspectable per book while retaining a complete ordinal distribution.

The empirical pattern is equally specific. Uniform pooling already supplies a strong heterogeneous baseline. Projection changes only nine main-set absolute errors relative to that pool (seven improve, two worsen), yet this is enough for a positive paired QWK interval and halves the frequency of errors above one level. Gains over length survive covariate control and counter-length comparisons, while the Qwen rubric baseline reveals substantial platform sensitivity despite explicit instructions not to use length alone. Capability routing further shows that illustration and pinyin criteria benefit from evidence unavailable to a text-only global model.

Why zero-parameter pooling rather than learned fusion.

Zero-parameter pooling is a design requirement, not a convenience. A learned fusion layer must be fitted on labelled books; because source platforms differ sharply in layout, length and illustration conventions, such a layer transfers its fitting-platform label distribution to new platforms, and its per-item weights are not inspectable. RA-SCOPE consumes exactly one quantity from held-out data—a label-free distance rank—so Eq. 7 can be re-executed and audited item by item. Where a fitted alternative is affordable it does raise average fit: the stacked fusion of Section VI-C reduces MAE by .142 relative to the contract. Our claim concerns decision guarantees and auditability under platform shift, not the best average fit on a labelled target collection.

IX. LIMITATIONS AND ETHICAL CONSIDERATIONS

Observed gold labels contain four L5 books and no L6 books, and the 28-book robustness collection is small; the results therefore support L1–L4 use most directly, with L5 descriptive and L6 untested. Material grades are not learner-comprehension outcomes. Architecture selection had access to the adjudicated collections, so a prospectively untouched, upper-level sample remains necessary for deployment-level validation. The conformal interpretation of the support rank requires exchangeable calibration and target representations, and platform shift can violate that assumption; a high support rank does not imply prediction correctness, and two supported anchors can share an error. StoryWeaver dominates the main collection, while the smallest source slice has eight books. The Qwen experiment is a reproducible 3B zero-shot rubric baseline, not a claim about all larger or proprietary LLMs. OCR, page extraction and DINOv2 style sensitivity can introduce source-specific error, and comparable geometric evidence is still needed for a dedicated layout specialist.

The study uses books and adult professional annotations, with no learner records or children’s personal data. Annotators are represented by role identifiers, and independent ratings precede separate adjudication. Predictions are intended to assist human material review: conflict, no-support and upper-band cases remain visible rather than being presented as certified decisions. Restricted book text and illustrations remain with their source providers, and users must preserve applicable attribution and licence conditions. The Standard may privilege

standard written Mandarin and conventional print layouts; dialectal, culturally specific, accessibility-relevant and visually driven materials require human validation.

X. CONCLUSION

We presented a support-backed full-consensus projection for ordinal picturebook grading. It combines four monotone experts without fitted fusion parameters, attaches an explicit support state to every item, and uniquely preserves every activated anchor-consensus grade. Across source-only out-of-fold, main and separate robustness evaluations it improves the strongest individual experts and the uniform pool while remaining stable across support and margin settings. The result is a compact inference architecture whose accuracy, decision semantics and evidence boundary can all be audited, and whose states translate directly into human review triage.

ACKNOWLEDGMENT

This work was supported by the National College Students’ Innovation and Entrepreneurship Training Program (No. 202610032033). The reproducibility artifact—the projection module, its invariant tests, every per-book prediction, the reported tables and the figure sources—is available at <https://github.com/yyting2006-ai/RA-SCOPE-artifact>.

REFERENCES

- [1] Center for Language Education and Cooperation, “Chinese graded readers standards for international Chinese language education,” Ministry of Education of the People’s Republic of China, Tech. Rep., 2024, international Chinese Education Chinese Reading Grading Standard; Chinese–English edition.
- [2] M. Xia, E. Kochmar, and T. Briscoe, “Text readability assessment for second language learners,” in *Proc. 11th Workshop on Innovative Use of NLP for Building Educational Applications (BEA)*, 2016, pp. 12–22.
- [3] T. Deutsch, M. Jasbi, and S. Shieber, “Linguistic features for readability assessment,” in *Proc. 15th Workshop on Innovative Use of NLP for Building Educational Applications (BEA)*, 2020, pp. 1–17.
- [4] B. W. Lee, Y. S. Jang, and J. Lee, “Pushing on text readability assessment: A transformer meets handcrafted linguistic features,” in *Proc. 2021 Conf. Empirical Methods in Natural Language Processing (EMNLP)*, 2021, pp. 10 669–10 686.
- [5] J. Zeng, X. Yu, X. Tong, and W. Xiao, “PromptARA: Improving deep representation in hybrid automatic readability assessment with prompt and orthogonal projection,” in *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023, pp. 15 360–15 371.
- [6] J. Zeng, X. Tong, X. Yu, W. Xiao, and Q. Huang, “InterpretARA: Enhancing hybrid automatic readability assessment with linguistic feature interpreter and contrastive learning,” in *Proc. AAAI Conf. Artificial Intelligence*, vol. 38, no. 17, 2024, pp. 19 497–19 505.
- [7] J. Zeng, X. Yu, X. Tong, and W. Xiao, “Self-supervised collaborative information bottleneck for text readability assessment,” in *Proc. AAAI Conf. Artificial Intelligence*, vol. 39, no. 24, 2025, pp. 25 814–25 822.
- [8] J. Zeng, Y. Xie, X. Yu, J. Lee, and D.-X. Zhou, “Enhancing automatic readability assessment with pre-training and soft labels for ordinal regression,” in *Findings of the Association for Computational Linguistics: EMNLP 2022*, 2022, pp. 4557–4568.
- [9] H. H. Lim and J. Lee, “Improving readability assessment with ordinal log-loss,” in *Proc. 19th Workshop on Innovative Use of NLP for Building Educational Applications (BEA)*, 2024, pp. 343–350.
- [10] W. Cao, V. Mirjalili, and S. Raschka, “Rank consistent ordinal regression for neural networks with application to age estimation,” *Pattern Recognition Letters*, vol. 140, pp. 325–331, 2020.
- [11] J. Li, J. Qiu, and C. Lu, “Quantification of textual linguistic features and level identification based on the Chinese graded readers standards for international Chinese language education,” *Journal of International Chinese Teaching*, vol. 2025, no. 3, pp. 47–56, 63, 2025.

- [12] X. Yang, J. Yang, and X. Li, “Chinese automatic readability assessment using adaptive pre-training and linguistic feature fusion,” in *Proc. 31st Int. Conf. Computational Linguistics (COLING)*, 2025, pp. 9013–9024.
- [13] W. Li, R. Sun, T. Zhang, and Y. Wu, “Going beyond passages: Readability assessment for book-level long texts,” in *Proc. 23rd Chinese National Conf. Computational Linguistics (CCL)*, 2024, pp. 1298–1309.
- [14] M. Pérez-Ortiz, P. A. Gutiérrez, and C. Hervás-Martínez, “Projection-based ensemble learning for ordinal regression,” *IEEE Transactions on Cybernetics*, vol. 44, no. 5, pp. 681–694, 2014.
- [15] R. E. Barlow, D. J. Bartholomew, J. M. Bremner, and H. D. Brunk, *Statistical Inference under Order Restrictions: The Theory and Application of Isotonic Regression*. New York, NY, USA: Wiley, 1972.
- [16] Y. Sun, Y. Ming, X. Zhu, and Y. Li, “Out-of-distribution detection with deep nearest neighbors,” in *Proc. 39th Int. Conf. Machine Learning (ICML)*, ser. Proceedings of Machine Learning Research, vol. 162, 2022, pp. 20 827–20 840.
- [17] S. Bates, E. Candès, L. Lei, Y. Romano, and M. Sesia, “Testing for outliers with conformal p-values,” *The Annals of Statistics*, vol. 51, no. 1, pp. 149–178, 2023.
- [18] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, “On calibration of modern neural networks,” in *Proc. 34th Int. Conf. Machine Learning (ICML)*, ser. Proceedings of Machine Learning Research, vol. 70, 2017, pp. 1321–1330.
- [19] Y. Ovadia, E. Fertig, J. Ren, Z. Nado, D. Sculley, S. Nowozin, J. Dillon, B. Lakshminarayanan, and J. Snoek, “Can you trust your model’s uncertainty? evaluating predictive uncertainty under dataset shift,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019.
- [20] Y. Geifman and R. El-Yaniv, “SelectiveNet: A deep neural network with an integrated reject option,” in *Proc. 36th Int. Conf. Machine Learning (ICML)*, ser. Proceedings of Machine Learning Research, vol. 97, 2019, pp. 2151–2159.
- [21] J. Xin, R. Tang, Y. Yu, and J. Lin, “The art of abstention: Selective prediction and error regularization for natural language processing,” in *Proc. 59th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2021, pp. 1040–1051.
- [22] Ministry of Education of the People’s Republic of China and State Language Commission, “Chinese proficiency grading standards for international Chinese language education (GF 0025–2021),” Ministry of Education of the People’s Republic of China, Tech. Rep., 2021.
- [23] Pratham Books, “StoryWeaver open content,” Online catalogue, 2026, accessed 2026-08-25.
- [24] Global African Storybook Project, “Global African storybook project repository,” GitHub repository, 2026, snapshot audited 2026-08-25.
- [25] Global Digital Library, “Global digital library content API,” API documentation, 2026, snapshot audited 2026-08-25.
- [26] A. F. Hayes and K. Krippendorff, “Answering the call for a standard reliability measure for coding data,” *Communication Methods and Measures*, vol. 1, no. 1, pp. 77–89, 2007.
- [27] G. Salton and C. Buckley, “Term-weighting approaches in automatic text retrieval,” *Information Processing and Management*, vol. 24, no. 5, pp. 513–523, 1988.
- [28] Y. Cui, W. Che, T. Liu, B. Qin, S. Wang, and G. Hu, “Revisiting pre-trained models for Chinese natural language processing,” in *Findings of the Association for Computational Linguistics: EMNLP 2020*, 2020, pp. 657–668.
- [29] S. Boyd and L. Vandenberghe, *Convex Optimization*. Cambridge, U.K.: Cambridge University Press, 2004.
- [30] Qwen Team, A. Yang, B. Yang, B. Zhang, B. Hui, B. Zheng, B. Yu, C. Li, D. Liu, F. Huang, H. Wei, H. Lin, J. Yang, J. Tu, J. Zhang, J. Yang, J. Zhou, J. Lin, K. Dang *et al.*, “Qwen2.5 technical report,” arXiv preprint arXiv:2412.15115, 2024.
- [31] J. Cohen, “Weighted kappa: Nominal scale agreement with provision for scaled disagreement or partial credit,” *Psychological Bulletin*, vol. 70, no. 4, pp. 213–220, 1968.
- [32] B. Efron and R. J. Tibshirani, *An Introduction to the Bootstrap*. New York, NY, USA: Chapman and Hall/CRC, 1993.
- [33] M. Oquab, T. Darcet, T. Moutakanni, H. V. Vo, M. Szafraniec, V. Khalidov, P. Fernandez, D. Haziza, F. Massa, A. El-Nouby, M. Assran, N. Ballas, W. Galuba, R. Howes, P.-Y. Huang, S.-W. Li, I. Misra, M. Rabbat, V. Sharma *et al.*, “DINOv2: Learning robust visual features without supervision,” *Transactions on Machine Learning Research*, 2024.